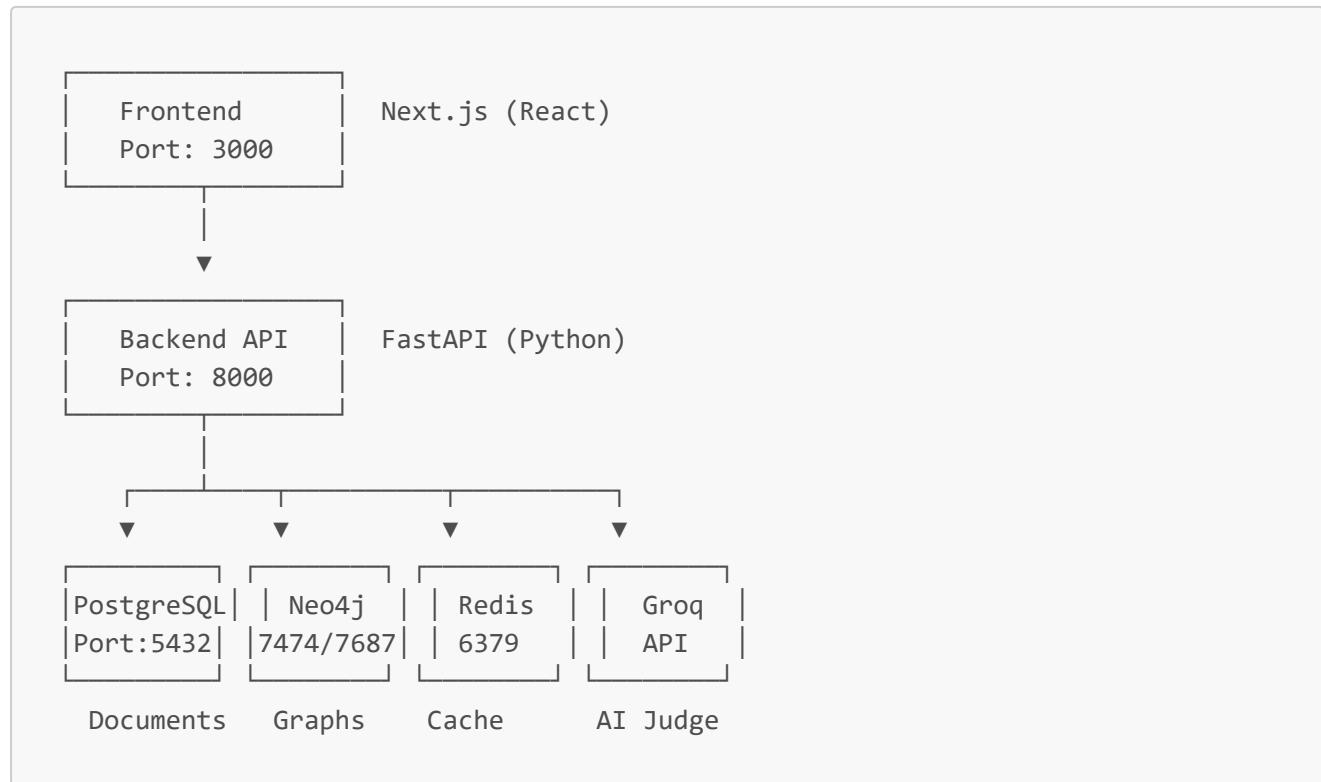


💡 AI Judge - Hybrid Database Setup

Architecture Overview



📦 What Each Database Does

PostgreSQL (Primary Storage)

- **Purpose:** Store case documents, user data, arguments
- **Data:** Text content, file metadata, verdicts
- **Why:** ACID compliance, reliable persistence, easy queries

Neo4j (Relationship Graph)

- **Purpose:** Track relationships between cases, arguments, precedents
- **Data:** Case similarities, argument patterns, winning strategies
- **Why:** Fast relationship traversal, pattern discovery, analytics

Redis (Cache Layer)

- **Purpose:** Cache verdicts, API responses, session data
- **Data:** Temporary cached results
- **Why:** Reduce AI API calls, faster response times, cost savings

🛠️ Setup Options

Option 1: Docker Compose (Recommended) 🤖

Fastest way to get everything running:

```
# 1. Copy environment file  
cp backend/.env.example backend/.env  
  
# 2. Add your Groq API key to backend/.env  
# GROQ_API_KEY=your_actual_key_here  
  
# 3. Start all services  
docker-compose up -d  
  
# 4. Wait for services to be healthy (30-60 seconds)  
docker-compose ps  
  
# 5. Access the services:  
# - Frontend: http://localhost:3000  
# - Backend API: http://localhost:8000  
# - Neo4j Browser: http://localhost:7474 (neo4j/password123)  
# - PostgreSQL: localhost:5432 (postgres/postgres123)  
# - Redis: localhost:6379
```

Stop services:

```
docker-compose down
```

View logs:

```
docker-compose logs -f backend
```

Option 2: Local Development (Manual)

If you want to run services individually:

1. Start Databases (Docker)

```
# PostgreSQL  
docker run -d --name ai-judge-postgres \  
-e POSTGRES_DB=ai_judge \  
-e POSTGRES_USER=postgres \  
-e POSTGRES_PASSWORD=postgres123 \  
-p 5432:5432 \  
postgres:15-alpine  
  
# Neo4j
```

```
docker run -d --name ai-judge-neo4j \
-e NEO4J_AUTH=neo4j/password123 \
-p 7474:7474 -p 7687:7687 \
neo4j:5.14-community

# Redis
docker run -d --name ai-judge-redis \
-p 6379:6379 \
redis:7-alpine
```

2. Setup Backend

```
cd backend

# Create virtual environment
python -m venv venv
.\venv\Scripts\Activate.ps1 # Windows
# source venv/bin/activate # Mac/Linux

# Install dependencies
pip install -r requirements.txt

# Create .env file
cp .env.example .env
# Edit .env and add your GROQ_API_KEY

# Update .env with local database URLs:
# DATABASE_URL=postgresql://postgres:postgres123@localhost:5432/ai_judge
# NEO4J_URI=bolt://localhost:7687
# NEO4J_USER=neo4j
# NEO4J_PASSWORD=password123
# REDIS_URL=redis://localhost:6379

# Run backend (use hybrid version)
python app_hybrid.py
```

3. Setup Frontend

```
cd frontend

# Install dependencies
npm install

# Create .env.local
echo "NEXT_PUBLIC_API_URL=http://localhost:8000" > .env.local
```

```
# Run development server  
npm run dev
```

Database Connections

PostgreSQL

```
# Connect via psql  
psql -h localhost -U postgres -d ai_judge  
# Password: postgres123  
  
# View tables  
\dt  
  
# Query cases  
SELECT * FROM cases;
```

Neo4j

```
# Open browser  
http://localhost:7474  
  
# Login: neo4j / password123  
  
# Sample queries:  
MATCH (c:Case) RETURN c LIMIT 10;  
MATCH (c:Case)-[r:SIMILAR_TO]-(similar:Case) RETURN c, r, similar;  
MATCH (c:Case)-[:HAS_VERDICT]->(v:Verdict) RETURN v.winning_side, count(*) as wins;
```

Redis

```
# Connect via redis-cli  
docker exec -it ai-judge-redis redis-cli  
  
# View all keys  
KEYS *  
# Get cached verdict  
GET verdict:case123
```

Testing the System

1. Create a Case

```
curl -X POST http://localhost:8000/api/case/create \
-H "Content-Type: application/json" \
-d '{"case_id": "test-case-001"}'
```

2. Submit Arguments

```
# Side A
curl -X POST http://localhost:8000/api/case/test-case-001/argument \
-H "Content-Type: application/json" \
-d '{
  "case_id": "test-case-001",
  "side": "A",
  "text": "Plaintiff argues breach of contract..."
}'

# Side B
curl -X POST http://localhost:8000/api/case/test-case-001/argument \
-H "Content-Type: application/json" \
-d '{
  "case_id": "test-case-001",
  "side": "B",
  "text": "Defendant argues force majeure..."
}'
```

3. Generate Verdict

```
curl -X POST http://localhost:8000/api/case/test-case-001/verdict
```

4. Check Statistics

```
curl http://localhost:8000/api/statistics
```

🔍 API Endpoints

Endpoint	Method	Description
/	GET	API info
/health	GET	Health check (all databases)

Endpoint	Method	Description
/api/case/create	POST	Create new case
/api/case/{id}/upload	POST	Upload documents
/api/case/{id}/argument	POST	Submit argument
/api/case/{id}/verdict	POST	Generate verdict
/api/case/{id}/followup	POST	Submit follow-up (max 5)
/api/case/{id}	GET	Get case details
/api/case/{id}/similar	GET	Find similar cases (Neo4j)
/api/statistics	GET	System statistics

📈 Scaling Features

Caching Strategy

- Verdicts cached for 24 hours
- Case data cached for 1 hour
- Redis automatically evicts old entries
- Cache invalidation on updates

Database Indexing

- PostgreSQL: Indexed on case_id, created_at, status
- Neo4j: Indexed on case_id for fast lookups
- Optimized JOIN queries with proper foreign keys

Connection Pooling

- SQLAlchemy manages PostgreSQL connections
- Neo4j driver uses connection pooling
- Redis connection pooling built-in

🔧 Troubleshooting

"Cannot connect to PostgreSQL"

```
# Check if running
docker ps | grep postgres

# Restart
docker restart ai-judge-postgres
```

```
# Check logs  
docker logs ai-judge-postgres
```

"Neo4j not available"

This is OK! Neo4j is optional. The system works without it, you just won't get:

- Case similarity matching
- Argument pattern analysis
- Relationship graphs

"Redis not available"

This is OK too! Redis is optional. Without it:

- No caching (slower responses)
- More API calls to Groq (higher cost)
- But everything still works

"Import errors in Python"

```
# Make sure you're in virtual environment  
.venv\Scripts\Activate.ps1  
  
# Reinstall requirements  
pip install -r requirements.txt --upgrade
```

🚀 Production Deployment

Database Options

Free Tier:

- PostgreSQL: Supabase (500MB free)
- Neo4j: Aura Free (200k nodes)
- Redis: Upstash (10k requests/day free)

Paid:

- PostgreSQL: AWS RDS, Google Cloud SQL
- Neo4j: Aura Professional
- Redis: Redis Enterprise Cloud

Environment Variables (Production)

```
# backend/.env
DATABASE_URL=postgresql://user:pass@host:5432/ai_judge
NEO4J_URI=neo4j+s://xxxxxx.databases.neo4j.io
NEO4J_USER=neo4j
NEO4J_PASSWORD=<secure-password>
REDIS_URL=redis://default:password@host:12345
GROQ_API_KEY=<your-key>
```

Additional Resources

- PostgreSQL Docs: <https://www.postgresql.org/docs/>
- Neo4j Docs: <https://neo4j.com/docs/>
- Redis Docs: <https://redis.io/documentation>
- FastAPI Docs: <https://fastapi.tiangolo.com/>
- SQLAlchemy Docs: <https://docs.sqlalchemy.org/>

Benefits of Hybrid Approach

Feature	Before (In-Memory)	After (Hybrid)
Persistence	 Lost on restart	 Permanent
Scalability	 Single server	 Horizontal scaling
Relationships	 Manual	 Graph queries
Performance	 Fast but limited	 Fast + cached
Analytics	 None	 Pattern discovery
Cost	 \$0	 \$0-50/month

You now have a production-ready, scalable AI Judge system!  