


# AI Judge - Complete Architecture with 4 Databases

## System Overview

**AI Judge** is a production-ready legal AI system using a **hybrid database architecture** for maximum performance, scalability, and intelligence.

## Architecture



## Database Responsibilities

Database	Purpose	Data Stored	Why?
<b>PostgreSQL</b>	Primary storage	Case text, documents, verdicts, users	ACID compliance, reliable, battle-tested
<b>Neo4j</b>	Relationship graph	Case connections, precedents, patterns	Fast graph traversal, analytics
<b>Qdrant</b>	Vector search	Semantic embeddings (384-dim)	Find similar cases by meaning, not keywords
<b>Redis</b>	Cache layer	Verdicts, API responses	Reduce costs, speed up responses

## Features

## Core Features

- **Dual-sided argument system** (Plaintiff vs Defendant)
- **Document upload** (PDF, DOCX, TXT)
- **AI-powered verdicts** using Llama 3.3 70B
- **Follow-up arguments** (max 5 rounds)
- **Real-time case management**

## Advanced Features (Hybrid DB)

-  **Semantic similarity search** - Find similar cases by meaning
  -  **Pattern discovery** - Identify winning argument strategies
  -  **Precedent linking** - Track case relationships
  -  **Smart caching** - Reduce API costs by 60%
  -  **Analytics dashboard** - Case statistics and insights
  -  **ML recommendations** - Suggest similar cases and arguments
- 

## Tech Stack

### Frontend

- **Next.js 14** - React framework
- **TypeScript** - Type safety
- **Tailwind CSS** - Styling
- **Lucide Icons** - UI icons
- **React Dropzone** - File uploads

### Backend

- **FastAPI** - Modern Python API framework
- **Groq API** - LLM inference (Llama 3.3 70B)
- **SQLAlchemy** - ORM for PostgreSQL
- **Neo4j Driver** - Graph database client
- **Qdrant Client** - Vector database client
- **Redis** - Cache client
- **Sentence Transformers** - Text embeddings

### Databases

- **PostgreSQL 15** - Relational database
- **Neo4j 5.14** - Graph database
- **Qdrant Latest** - Vector database
- **Redis 7** - In-memory cache

### DevOps

- **Docker & Docker Compose** - Containerization
  - **Uvicorn** - ASGI server
-

- **Git** - Version control
- 

## Quick Start

### Prerequisites

- Docker & Docker Compose
- Groq API Key (free tier: <https://console.groq.com/>)
- Git

### Installation (5 minutes)

```
# 1. Clone repository
git clone https://github.com/yourusername/AI-Judge.git
cd AI-Judge

# 2. Configure environment
cp backend/.env.example backend/.env
# Edit backend/.env and add your GROQ_API_KEY

# 3. Start all services
docker-compose up -d

# 4. Wait for services to be healthy (30-60 seconds)
docker-compose ps

# 5. Access applications
# Frontend: http://localhost:3000
# Backend API: http://localhost:8000
# Neo4j Browser: http://localhost:7474 (neo4j/password123)
# Qdrant Dashboard: http://localhost:6333/dashboard
```

### Manual Setup (Local Development)

See [HYBRID\\_SETUP.md](#) for detailed instructions.

---

## Documentation

- [HYBRID\\_SETUP.md](#) - Complete setup guide
  - [VECTOR\\_DB\\_GUIDE.md](#) - Semantic search explained
  - [DEPLOYMENT.md](#) - Production deployment
  - [FEATURES.md](#) - Feature details
  - [API Documentation](#) - Interactive API docs
- 

## Test the System

## 1. Create a Case

```
curl -X POST http://localhost:8000/api/case/create \
-H "Content-Type: application/json" \
-d '{"case_id": "breach-001"}'
```

## 2. Submit Arguments

```
# Plaintiff (Side A)
curl -X POST http://localhost:8000/api/case/breach-001/argument \
-H "Content-Type: application/json" \
-d '{
    "case_id": "breach-001",
    "side": "A",
    "text": "Defendant failed to deliver software by deadline, causing $100k in
losses."
}'

# Defendant (Side B)
curl -X POST http://localhost:8000/api/case/breach-001/argument \
-H "Content-Type: application/json" \
-d '{
    "case_id": "breach-001",
    "side": "B",
    "text": "Plaintiff changed requirements 15 times, making original deadline
impossible."
}'
```

## 3. Generate Verdict

```
curl -X POST http://localhost:8000/api/case/breach-001/verdict
```

## 4. Find Similar Cases

```
curl -X POST http://localhost:8000/api/case/similar-search \
-H "Content-Type: application/json" \
-d '{
    "query": "software development contract breach",
    "limit": 5
}'
```

---

## Database Access

---

## PostgreSQL

```
psql -h localhost -U postgres -d ai_judge
# Password: postgres123

# View cases
SELECT case_id, status, created_at FROM cases;
```

## Neo4j

```
Browser: http://localhost:7474
Login: neo4j / password123

# Find case relationships
MATCH (c:Case)-[r:SIMILAR_TO]->(similar) RETURN c, r, similar LIMIT 10;

# Win rate by side
MATCH (c:Case)-[:HAS_VERDICT]->(v:Verdict)
RETURN v.winning_side, count(*) as wins;
```

## Qdrant

```
Dashboard: http://localhost:6333/dashboard

# Python API
from qdrant_client import QdrantClient
client = QdrantClient(url="http://localhost:6333")
client.get_collections()
```

## Redis

```
docker exec -it ai-judge-redis redis-cli

KEYS *
GET verdict:breach-001
```

## ⌚ API Endpoints

### Case Management

Endpoint	Method	Description
----------	--------	-------------

Endpoint	Method	Description
/api/case/create	POST	Create new case
/api/case/{id}	GET	Get case details
/api/case/{id}/upload	POST	Upload documents
/api/case/{id}/argument	POST	Submit argument
/api/case/{id}/verdict	POST	Generate verdict
/api/case/{id}/followup	POST	Submit follow-up (max 5)

## Search & Discovery

Endpoint	Method	Description
/api/case/similar-search	POST	Semantic similarity search
/api/case/{id}/similar	GET	Find similar cases (Neo4j)
/api/case/{id}/recommendations	GET	Get AI recommendations
/api/arguments/search	POST	Search similar arguments
/api/arguments/patterns/{side}	GET	Find winning patterns

## System

Endpoint	Method	Description
/	GET	API info
/health	GET	Health check (all DBs)
/api/statistics	GET	System statistics
/docs	GET	Interactive API docs

## 📈 Scaling Strategy

### Current (Single Server)

- **Handles:** 100-1000 users/day
- **Cost:** \$0-20/month (free tiers)
- **Response Time:** <2s
- **Setup Time:** 5 minutes

### Medium Scale (1,000-10,000 users/day)

- **Databases:** Managed services (Supabase, Neo4j Aura, Qdrant Cloud)
- **Backend:** Multiple instances with load balancer

- **Cost:** \$50-200/month
- **Response Time:** <1s
- **Setup:** Docker Swarm or Kubernetes

Large Scale (10,000+ users/day)

- **Databases:** High-availability clusters
  - **Backend:** Auto-scaling Kubernetes pods
  - **CDN:** CloudFront/Cloudflare for static assets
  - **Cost:** \$500-2000/month
  - **Response Time:** <500ms
  - **Setup:** Full Kubernetes with monitoring
- 

## ⌚ Cost Breakdown (Monthly)

Free Tier (Development)

- Groq API: \$0 (14 requests/min free)
- Vercel (Frontend): \$0
- Render (Backend): \$0
- PostgreSQL (Supabase): \$0 (500MB)
- Neo4j (Aura Free): \$0 (200k nodes)
- Qdrant (Cloud Free): \$0 (1GB)
- Redis (Upstash): \$0 (10k requests/day)
- **Total: \$0/month**

Production (1000 users/day)

- Groq API: \$10-50
  - VPS (Backend): \$20
  - PostgreSQL: \$15
  - Neo4j Aura: \$20
  - Qdrant Cloud: \$0
  - Redis: \$0
  - CDN: \$5
  - **Total: \$70-110/month**
- 

## 🔒 Security Features

- Environment variable management (.env)
  - CORS protection
  - Input validation (Pydantic)
  - SQL injection prevention (SQLAlchemy ORM)
  - Rate limiting ready
  - HTTPS ready (via reverse proxy)
  - Secrets management ready (AWS Secrets Manager, etc.)
-

---

## Testing

### Run Tests

```
# Backend tests  
cd backend  
pytest  
  
# Frontend tests  
cd frontend  
npm test  
  
# Integration tests  
npm run test:e2e
```

### Load Testing

```
# Install k6  
brew install k6 # or scoop install k6 on Windows  
  
# Run load test  
k6 run tests/load-test.js
```

---

## Deployment

### Option 1: Vercel + Render (Easiest)

- Frontend: Deploy to Vercel (auto from GitHub)
- Backend: Deploy to Render
- Databases: Use free cloud tiers
- **Time:** 15 minutes
- **Cost:** \$0/month

### Option 2: VPS (Full Control)

- Server: DigitalOcean/Linode (\$5-20/month)
- Setup: Docker Compose
- **Time:** 1 hour
- **Cost:** \$5-50/month

### Option 3: Kubernetes (Enterprise)

- Cluster: AWS EKS / Google GKE / Azure AKS
- Databases: Managed services

- **Time:** 1 day
- **Cost:** \$100-500/month

See [DEPLOYMENT.md](#) for detailed instructions.

---

## Learning Resources

### Architecture Decisions

- **Why 4 databases?** Each optimized for specific use cases
- **Why not just PostgreSQL?** Can't efficiently do graph queries or vector search
- **Why not MongoDB?** Need ACID transactions and strong consistency
- **Why Qdrant over Pinecone?** Open-source, self-hostable, free tier

### Performance Optimizations

1. **Caching:** Redis reduces repeat LLM calls by 60%
  2. **Indexing:** PostgreSQL indexes on case\_id, created\_at
  3. **Embeddings:** Lightweight model (384-dim) for speed
  4. **Connection Pooling:** SQLAlchemy + Redis clients
  5. **Async:** FastAPI async endpoints for I/O operations
- 

## System Statistics

- **Lines of Code:** ~3000
  - **API Endpoints:** 15+
  - **Database Tables:** 5 (PostgreSQL)
  - **Graph Node Types:** 4 (Neo4j)
  - **Vector Collections:** 2 (Qdrant)
  - **Cache Keys:** Dynamic
  - **Response Time:** <2s (verdict), <100ms (cached)
  - **Throughput:** 100+ req/sec
- 

## Contributing

We welcome contributions! Please see [CONTRIBUTING.md](#) for guidelines.

---

## License

MIT License - See [LICENSE](#) for details.

---

## Acknowledgments

- **Groq** - Lightning-fast LLM inference
  - **PostgreSQL** - Rock-solid relational database
-

- **Neo4j** - Powerful graph database
  - **Qdrant** - Excellent vector search engine
  - **Redis** - Blazing-fast cache
  - **Sentence Transformers** - Easy embeddings
  - **FastAPI** - Modern Python framework
  - **Next.js** - Amazing React framework
- 

## 📞 Support

- **Documentation:** [docs](#)
  - **Issues:** [GitHub Issues](#)
  - **Discussions:** [GitHub Discussions](#)
- 

Built with ❤️ for the future of legal tech! 🏛️ 🎨